

## Analyze A/B Test Results

This project will assure you have mastered the subjects covered in the statistics lessons. The hope is to have this project be as comprehensive of these topics as possible. Good luck!

## Table of Contents

- [Introduction](#)
- [Part I - Probability](#)
- [Part II - A/B Test](#)
- [Part III - Regression](#)

### Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these

For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

**As you work through this notebook, follow along in the classroom and answer the corresponding quiz questions associated with each question.** The labels for each classroom concept are provided for each question. This will assure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the criteria. As a final check, assure you meet all the criteria on the [RUBRIC](#).

### Part I - Probability

To get started, let's import our libraries.

```
In [1]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

1. Now, read in the `ab_data.csv` data. Store it in `df`. Use your dataframe to answer the questions in Quiz 1 of the classroom.

a. Read in the dataset and take a look at the top few rows here:

```
In [2]: df = pd.read_csv('./ab_data.csv')

In [3]: df.head()
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

b. Use the below cell to find the number of rows in the dataset.

```
In [4]: df.shape[0]
Out[4]: 294478

In [5]: df['user_id'].nunique()
Out[5]: 290584

In [6]: df['converted'].mean() * 100
Out[6]: 11.9691935560551

In [7]: df.groupby(["group", "landing_page"]).size()
Out[7]:
group  landing_page  1928
control  old_page      145274
         new_page      145311
treatment  old_page      1965
           old_page      1965
dtype: int64

In [8]: df.query("group == 'treatment' and landing_page == 'old_page'").shape[0] + df.query("group == 'control'
and landing_page == 'new_page'").shape[0]
Out[8]: 3893

f. Do any of the rows have missing values?
```

```
In [9]: df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  --
0   user_id      294478 non-null  int64
1   timestamp    294478 non-null  object
2   group        294478 non-null  object
3   landing_page  294478 non-null  object
4   converted     294478 non-null  int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB
```

2. For the rows where `treatment` is not aligned with `new_page` or `control` is not aligned with `old_page`, we cannot be sure if this row truly received the new or old page. Use **Quiz 2** in the classroom to provide how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in `df2`.

```
In [10]: df2 = df.query("group == 'control' and landing_page == 'old_page'")
df2 = df2.append(df.query("group == 'treatment' and landing_page == 'new_page'"))
df2.head()
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1
5	936923	2017-01-10 15:20:49.083499	control	old_page	0
7	719014	2017-01-17 01:48:29.539573	control	old_page	0

In [11]: # Double Check all of the correct rows were removed - this should be 0
df2[(df2["group"] == 'treatment') == (df2["landing\_page"] == 'new\_page')] == False].shape[0]
Out[11]: 0

3. Use `df2` and the cells below to answer questions for Quiz3 in the classroom.

a. How many unique `user_ids` are in `df2`?

```
In [12]: df2['user_id'].nunique()
Out[12]: 290584

b. There is one user_id repeated in df2. What is it?
```

```
In [13]: df2['user_id'].duplicated().sum()
Out[13]: 1

df2[df2['user_id'].duplicated()]
```

	user_id	timestamp	group	landing_page	converted
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

c. What is the row information for the repeat `user_id`?

```
In [15]: df2[df2['user_id'] == 773192]
Out[15]:
   user_id  timestamp      group  landing_page  converted
1899  773192  2017-01-09 05:37:58.781806  treatment    new_page      0
2893  773192  2017-01-14 02:55:59.590927  treatment    new_page      0

d. Remove one of the rows with a duplicate user_id, but keep your dataframe as df2.
```

```
In [16]: df2 = df2.drop_duplicates('user_id')
In [17]: df2['user_id'].duplicated().sum()
Out[17]: 0

4. Use df2 in the below cells to answer the quiz questions related to Quiz 4 in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?


```

```
In [18]: df2.converted.mean()
Out[18]: 0.11959708724499628

P(converted) = 0.1196

b. Given that an individual was in the control group, what is the probability they converted?
```

```
In [19]: df2.query("group == 'control'")['converted'].mean()
Out[19]: 0.1203863045004612

P(converted | control) = 0.1204

c. Given that an individual was in the treatment group, what is the probability they converted?
```

```
In [20]: df2.query("group == 'treatment'")['converted'].mean()
Out[20]: 0.11880860551510564

P(converted | treatment) = 0.1188

d. What is the probability that an individual received the new page?
```

```
In [21]: df2.query("landing_page == 'new_page'").shape[0] / df2.shape[0]
Out[21]: 0.5000619442226688

P(new page) = 0.5001
```

e. Consider your results from a. through d. above, and explain below whether you think there is sufficient evidence to say that the new treatment page leads to more conversions.

**Not sufficient.**

The new treatment page  $P(\text{converted} | \text{treatment}) = 0.1188$  leads to lower conversions rate than the old control page  $P(\text{converted} | \text{control}) = 0.1204$ , but the difference appears to be negligible and not sufficient evidence.

### Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better then the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of  $\mu_{\text{old}}$  and  $\mu_{\text{new}}$ , which are the converted rates for the old and new pages.

$H_0: \mu_{\text{new}} \leq \mu_{\text{old}}$

$H_1: \mu_{\text{new}} > \mu_{\text{old}}$

2. Assume under the null hypothesis,  $\mu_{\text{new}}$  and  $\mu_{\text{old}}$  both have "true" success rates equal to the converted success rate regardless of page - that is  $\mu_{\text{new}}$  and  $\mu_{\text{old}}$  are equal. Furthermore, assume they are equal to the converted rate in `ab_data.csv` to make the page.

Use a sample size for each page equal to the ones in `ab_data.csv`.

Perform the sampling distribution for the difference in converted between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use **Quiz 5** in the classroom to make sure you are on the right track.

a. What is the convert rate for  $\mu_{\text{new}}$  under the null?

```
In [22]: p_new = df2['converted'].mean()
Out[22]: 0.11959708724499628

b. What is the convert rate for  $\mu_{\text{old}}$  under the null?
```

```
In [23]: p_old = df2['converted'].mean()
p_old
Out[23]: 0.11959708724499628

c. What is  $\mu_{\text{new}}$ ?
```

```
In [24]: n_new = df2.query("group == 'treatment'").shape[0]
n_new
Out[24]: 145310

d. What is  $\mu_{\text{old}}$ ?
```

```
In [25]: n_old = df2.query("group == 'control'").shape[0]
n_old
Out[25]: 145274
```

e. Simulate  $\mu_{\text{new}}$  transactions with a convert rate of  $\mu_{\text{new}}$  under the null. Store these  $\mu_{\text{new}}$  1's and 0's in `new_page_converted`.

```
In [26]: new_page_converted = np.random.binomial(n_new, p_new)
new_page_converted
Out[26]: 17374

f. Simulate  $\mu_{\text{old}}$  transactions with a convert rate of  $\mu_{\text{old}}$  under the null. Store these  $\mu_{\text{old}}$  1's and 0's in old_page_converted.
```

```
In [27]: old_page_converted = np.random.binomial(n_old, p_old)
old_page_converted
Out[27]: 17319

g. Find  $\mu_{\text{new}} - \mu_{\text{old}}$  for your simulated values from part (e) and (f).
```

```
In [28]: (new_page_converted / n_new) - (old_page_converted / n_old)
Out[28]: 0.00034896579952159446
```

h. Simulate 10,000  $\mu_{\text{new}} - \mu_{\text{old}}$  values using this same process similarly to the one you calculated in parts a. through g. above. Store all 10,000 values in a numpy array called `p_diffs`.

```
In [29]: p_diffs = (np.random.binomial(n_new, p_new, 10000) / n_new) - (np.random.binomial(n_old, p_old, 10000) / n_old)
p_diffs = np.array(p_diffs)
Out[29]: array([-1.08968471e-03,  1.10616589e-03, -1.48880237e-03, ...,
        2.72999534e-04, -1.87827861e-04,  1.84192482e-05])

i. Plot a histogram of the p_diffs. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.
```

```
In [30]: plt.figure(figsize=(12,6))
plt.hist(p_diffs)
plt.title('Histogram of 10,000 simulated pages difference',fontsize=15)
plt.xlabel('nPages difference', fontsize=15)
plt.ylabel('Frequency', fontsize=15)
plt.savefig(x = obs_diffs, color = 'red')
Out[30]: Text(0, 0.5, 'Frequency')
```

j. What proportion of the `p_diffs` are greater than the actual difference observed in `ab_data.csv`?

```
In [31]: obs_diff = df2.query("group == 'treatment'")['converted'].mean() - df2.query("group == 'control'")['con
verted'].mean()
(p_diffs > obs_diff).mean()
Out[31]: 0.9029

In [32]: plt.figure(figsize=(12,6))
plt.hist(p_diffs)
plt.title('Histogram of 10,000 simulated pages difference',fontsize=15)
plt.xlabel('nPages difference', fontsize=15)
plt.ylabel('Frequency', fontsize=15)
plt.savefig(x = obs_diff, color = 'red')
Out[32]: <matplotlib.lines.Line2D at 0x286142bfb38>
```

k. In words, explain what you just computed in part j. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

- What is this value called in scientific studies? The **p-value** is the probability of obtaining a statistic as extreme or more extreme than the one observed in the experiment.
- What does this value mean in terms of whether or not there is a difference between the new and old pages? A large p-value in this case indicates that there is a slightly larger conversion rate for the new treatment page than the old control page.

l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer the the number of rows associated with the old page and new pages, respectively.

```
In [33]: import statsmodels.api as sm
convert_old = df2.query("group == 'control'")['converted'].sum()
convert_new = df2.query("group == 'treatment'")['converted'].sum()
n_old = df2.query("group == 'control'").shape[0]
n_new = df2.query("group == 'treatment'").shape[0]

m. Now use stats.proportions_ztest to compute your test statistic and p-value. Here is a helpful link on using the built in.
```

```
In [34]: z_score, p_value = sm.stats.proportions_ztest([convert_old, convert_new], [n_old, n_new], alternative=
'smaller')
results = Logit.fit()
Out[34]: (1.3109241984234394, 0.9050583127590245)
```

n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts j. and k.?

```
In [35]: from scipy.stats import norm
norm.cdf(z_score)
Out[35]: 0.9050583127590245

In [36]: norm.ppf(1-(0.05))
Out[36]: 1.6448536269514722
```

- What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? The z-score is 1.31 is inside our critical value = 1.64 and the p-value = 0.905 is still large. So the null hypothesis  $H_0$  that the conversion rate of the old pages is higher than the conversion rate of the new pages is not rejected.
- Do they agree with the findings in parts j. and k.? These values agree with the findings in parts j. and k.

### Part III - A regression approach

1. In this final part, you will see that the result you achieved in the previous A/B test can also be achieved by performing regression.

a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

**Logistic Regression**

b. The goal is to use `statsmodels` to fit the regression model you specified in part a. to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create a column for the intercept, and create a dummy variable column for which page each user received. Add an `intercept` column, as well as an `ab_page` column, which is 1 when an individual receives the treatment and 0 if control.

```
In [37]: df3 = df.copy()
df3.head()
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

```
In [38]: df3['intercept'] = 1
df3['ab_page'] = pd.get_dummies(df3['group'])['treatment']
df3
```

	user_id	timestamp	group	landing_page	converted	intercept	ab_page
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	1	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	1	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	1	1
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	1	1
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	1	0
...	...	...	...	...	...	...	...
294473	751197	2017-01-03 22:28:38.630509	control	old_page	0	1	0
294474	945152	2017-01-12 00:51:57.078372	control	old_page	0	1	0
294475	734608	2017-01-22 11:45:03.439544	control	old_page	0	1	0
294476	697314	2017-01-15 01:20:28.957438	control	old_page	0	1	0
294477	715931	2017-01-16 12:40:24.467417	treatment	new_page	0	1	1

294478 rows x 7 columns

c. Use `statsmodels` to import your regression model. Instantiate the model, and fit the model using the two columns you created in part b. to predict whether or not an individual converts.

```
In [39]: import statsmodels.api as sm
logit = sm.Logit(df3[['converted'],df3[['intercept','ab_page']]])
results = logit.fit()
Optimization terminated successfully.
Current function value: 0.366243
Iterations: 6

d. Provide the summary of your model below, and use it as necessary to answer the following questions.
```

```
In [40]: results.summary()
Out[40]:
Logit Regression Results
Dep. Variable:   converted   No. Observations:   294478
Model:            Logit      DF Residuals:    294476
Method:            MLE        DF Model:        1
Date:    Sun, 11 Jul 2021    Pseudo R-squ.:   7.093e-06
Time:            19:08:15    Log-Likelihood:  -1.0785e+05
converged:         True      LL-Null:         -1.0785e+05
Covariance Type: nonrobust    LLR P-value:      0.2161
```

	coef	std err	z	P> z	[0.025	0.975]
intercept	-1.9887	0.008	-248.297	0.000	-2.004	-1.973
ab_page	-0.0140	0.011	-1.237	0.216	-0.036	0.008

e. What is the p-value associated with `ab_page`? Why does it differ from the value you found in Part II?

Hint: What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in the Part II?

- What is the p-value associated with `ab_page`? 0.216
- Why does it differ from the value you found in Part II? The null and alternative hypotheses for the regression model are:  $H_0: \mu_{\text{new}} \leq \mu_{\text{old}}$   $H_1: \mu_{\text{new}} > \mu_{\text{old}}$  This only predicts the difference in the two values and while Part II predicts which page gets more conversions rate.

f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

- Other factors will help the regression model to get more accurate predictions of whether or not an individual converts.
- A disadvantage to adding additional terms into your regression model is that the model gets more complex and can cause bias in testing and interpreting.

g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives. You will need to read in the `countries.csv` dataset and merge together your datasets on the appropriate rows. [Here](#) are the docs for joining tables. Does it appear that country has an impact on conversion? Don't forget to create dummy variables for these country columns - Hint: You will need two columns for the three dummy variables. Provide the statistical output as well as a written response to answer this question.

```
In [41]: countries_df = pd.read_csv('./countries.csv')
countries_df.head()
```

	user_id	country
0	834778	UK
1	928468	US
2	822059	UK
3	711597	UK
4	710616	UK

```
In [42]: df_new = countries_df.set_index('user_id').join(df2.set_index('user_id'), how='inner')
df_new.head()
```

	country	timestamp	group	landing_page	converted
user_id					
834778	UK	2017-01-14 23:08:43.304998	control	old_page	0
928468	US	2017-01-23 14:44:16.387854	treatment	new_page	0
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	1
711597	UK	2017-01-22 03:14:24.763511	control	old_page	0
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	0
...	...	...	...	...	...
653118	US	2017-01-09 03:12:31.034796	control	old_page	0
878226	UK	2017-01-06 15:02:50.334962	control	old_page	0
799368	UK	2017-01-09 18:07:34.253935	control	old_page	0
655535	CA	2017-01-09 13:30:47.524512	treatment	new_page	0
934996	UK	2017-01-09 00:30:08.377677	control	old_page	0

290584 rows x 10 columns

```
In [45]: ## Fit Your Linear Model And Obtain the Results
logit = sm.Logit(df_new[['converted'],df_new[['intercept','US','UK']]])
results = logit.fit()
Optimization terminated successfully.
Current function value: 0.366116
Iterations: 6

Logit Regression Results
Dep. Variable:   converted   No. Observations:   290584
Model:            Logit      DF Residuals:    290581
Method:            MLE        DF Model:        2
Date:    Sun, 11 Jul 2021    Pseudo R-squ.:   1.521e-05
Time:            19:08:16    Log-Likelihood:  -1.0639e+05
converged:         True      LL-Null:         -1.0639e+05
Covariance Type: nonrobust    LLR P-value:      0.1994
```

	coef	std err	z	P> z	[0.025	0.975]
intercept	-1.9898	0.007	-292.314	0.000	-2.010	-1.983
US	-0.0408	0.027	-1.518	0.129	-0.093	0.012
UK	0.0099	0.013	0.748	0.456	-0.016	0.036

It doesn't appear that country a user life affects the conversion rate because all the p\_values are still larger than the Type I Error 0.05.

h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model. Provide the summary results, and your conclusions based on the results.

```
In [46]: df_new[['ab_page']] = pd.get_dummies(df_new[['group']])['treatment']
df_new
```

	country	timestamp	group	landing_page	converted	intercept	US	UK	CA	ab_page
user_id										
834778	UK	2017-01-14 23:08:43.304998	control	old_page	0	1	0	1	0	0
928468	US	2017-01-23 14:44:16.387854	treatment	new_page	0	1	0	0	1	1
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	1	1	0	1	0	1
711597	UK	2017-01-22 03:14:24.763511	control	old_page	0	1	0	1	0	0
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	0					